

```
In [1]: 1 from pyspark.sql import SparkSession
2 import pyspark.sql.functions as F
3 from pyspark.sql.types import *
4
5 spark = SparkSession\
6     .builder\
7     .appName("chapter-25-ML-preprocessing")\
8     .getOrCreate()
9
10 import os
11 SPARK_BOOK_DATA_PATH = os.environ['SPARK_BOOK_DATA_PATH']
```

```
In [41]: 1 from IPython.display import display
```

```
In [2]: 1 sales = spark.read.format("csv")\
2     .option("header", "true")\
3     .option("inferSchema", "true")\
4     .load(SPARK_BOOK_DATA_PATH + "/data/retail-data/by-day/*.csv")\
5     .coalesce(5)\
6     .where("Description IS NOT NULL")
```

```
In [3]: 1 sales.show(3, False)
```

```
+-----+-----+-----+-----+-----+-----+
+-----+
|InvoiceNo|StockCode|Description|Quantity|InvoiceDate|UnitPrice|CustomerID|Country|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|580538|23084|RABBIT NIGHT LIGHT|48|2011-12-05 08:38:00|1.79|14075.0|United Kingdom|
|580538|23077|DOUGHNUT LIP GLOSS|20|2011-12-05 08:38:00|1.25|14075.0|United Kingdom|
|580538|22906|12 MESSAGE CARDS WITH ENVELOPES|24|2011-12-05 08:38:00|1.65|14075.0|United Kingdom|
```

only showing top 3 rows

```
In [4]: 1 sales.count()
```

```
Out[4]: 540455
```

```
In [5]: 1 fakeIntDF = spark.read.parquet(SPARK_BOOK_DATA_PATH + "/data/simple-ml-integers")
```

```
In [6]: 1 fakeIntDF.show(5, False)
```

```
+----+----+----+
|int1|int2|int3|
+----+----+----+
|4    |5    |6    |
|1    |2    |3    |
|7    |8    |9    |
+----+----+----+
```

```
In [7]: 1 simpleDF = spark.read.json(SPARK_BOOK_DATA_PATH + "/data/simple-ml")
```

```
In [8]: 1 simpleDF.show(5, False)
```

```
+-----+-----+-----+-----+
|color|lab |value1|value2          |
+-----+-----+-----+-----+
|green|good|1      |14.386294994851129|
|blue |bad |8      |14.386294994851129|
|blue |bad |12     |14.386294994851129|
|green|good|15     |38.97187133755819 |
|green|good|12     |14.386294994851129|
+-----+-----+-----+-----+
only showing top 5 rows
```

```
In [9]: 1 scaledDF = spark.read.parquet(SPARK_BOOK_DATA_PATH + "/data/simple-ml-scaling")
```

```
1 scaleDF.show(5, False)
```

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import RFormula
4
5 supervised = RFormula(formula="lab ~ . + color:value1 + color:value2")
6 supervised.fit(simpleDF).transform(simpleDF).show(5, False)
```

```
In [12]: 1 # SQLTransformer - like spark.sql()
2
3 from pyspark.ml.feature import SQLTransformer
4
5 basicTransformation = SQLTransformer()\
6     .setStatement("""
7         SELECT sum(Quantity), count(*), CustomerID
8         FROM __THIS__
9         GROUP BY CustomerID
10     """)
11
12 basicTransformation.transform(sales).show(5, False)
```

```
+-----+-----+-----+
|sum(Quantity)|count(1)|CustomerID|
+-----+-----+-----+
|119          |62       |14452.0    |
|440          |143      |16916.0    |
|630          |72       |17633.0    |
|34           |6        |14768.0    |
|1542         |30       |13094.0    |
+-----+-----+-----+
only showing top 5 rows
```

```
In [13]: 1 sales.createOrReplaceTempView("sales")
```

```
In [14]: 1 spark.sql("select sum(Quantity), count(*), CustomerID from sales group by CustomerID").show(5, False)
```

sum(Quantity)	count(1)	CustomerID
119	62	14452.0
440	143	16916.0
630	72	17633.0
34	6	14768.0
1542	30	13094.0

only showing top 5 rows

```
In [15]: 1 # VectorAssembler - transformer to assemble columns into vector
2
3 from pyspark.ml.feature import VectorAssembler
4 va = VectorAssembler().setInputCols(["int1", "int2", "int3"])
5 va.transform(fakeIntDF).show(5, False)
```

int1	int2	int3	VectorAssembler_91f1042ca8cb__output
4	5	6	[4.0,5.0,6.0]
1	2	3	[1.0,2.0,3.0]
7	8	9	[7.0,8.0,9.0]

In []:

```
1
```

In [16]:

```
1 # COMMAND -----  
2  
3 contDF = spark.range(20).selectExpr("cast(id as double)")  
4 contDF.show(5, False)
```

```
+---+  
|id |  
+---+  
|0.0|  
|1.0|  
|2.0|  
|3.0|  
|4.0|  
+---+
```

only showing top 5 rows

```
In [17]: 1 # Bucketizer - transformer to split data into buckets
2
3 from pyspark.ml.feature import Bucketizer
4 bucketBorders = [-1.0, 5.0, 10.0, 250.0, 600.0]
5 bucketer = Bucketizer().setSplits(bucketBorders).setInputCol("id")
6 bucketer.transform(contDF).show(truncate=False)
```

```
+---+-----+
|id  |Bucketizer_bffc8656d38c__output|
+---+-----+
|0.0  |0.0                               |
|1.0  |0.0                               |
|2.0  |0.0                               |
|3.0  |0.0                               |
|4.0  |0.0                               |
|5.0  |1.0                               |
|6.0  |1.0                               |
|7.0  |1.0                               |
|8.0  |1.0                               |
|9.0  |1.0                               |
|10.0 |2.0                               |
|11.0 |2.0                               |
|12.0 |2.0                               |
|13.0 |2.0                               |
|14.0 |2.0                               |
|15.0 |2.0                               |
|16.0 |2.0                               |
|17.0 |2.0                               |
|18.0 |2.0                               |
|19.0 |2.0                               |
+---+-----+
```

```
In [18]: 1 # QuantileDiscretizer - transformer to partition data by Percentile
2 from pyspark.ml.feature import QuantileDiscretizer
3 bucketer = QuantileDiscretizer().setInputCol("id").setNumBuckets(5)
```

```
In [19]: 1 bucketer.fit(contDF).transform(contDF).show(truncate=False)
```

```
+-----+-----+
|id  |QuantileDiscretizer_4f6bd0bcc5c0__output|
+-----+-----+
|0.0  |0.0  |
|1.0  |0.0  |
|2.0  |0.0  |
|3.0  |1.0  |
|4.0  |1.0  |
|5.0  |1.0  |
|6.0  |1.0  |
|7.0  |2.0  |
|8.0  |2.0  |
|9.0  |2.0  |
|10.0 |2.0  |
|11.0 |3.0  |
|12.0 |3.0  |
|13.0 |3.0  |
|14.0 |3.0  |
|15.0 |4.0  |
|16.0 |4.0  |
|17.0 |4.0  |
|18.0 |4.0  |
|19.0 |4.0  |
+-----+-----+
```



```
In [20]: 1 # StandardScaler - transformer to normalize data
2
3 from pyspark.ml.feature import StandardScaler
4 sScaler = StandardScaler().setInputCol("features")
5 sScaler.fit(scaledDF).transform(scaledDF).show(5, False)
```

id	features	StandardScaler_3325a6ec77c0__output
0	[1.0, 0.1, -1.0]	[1.1952286093343936, 0.02337622911060922, -0.5976143046671968]
1	[2.0, 1.1, 1.0]	[2.390457218668787, 0.2571385202167014, 0.5976143046671968]
0	[1.0, 0.1, -1.0]	[1.1952286093343936, 0.02337622911060922, -0.5976143046671968]
1	[2.0, 1.1, 1.0]	[2.390457218668787, 0.2571385202167014, 0.5976143046671968]
1	[3.0, 10.1, 3.0]	[3.5856858280031805, 2.3609991401715313, 1.7928429140015902]

```
In [21]: 1 # COMMAND -----
2
3 from pyspark.ml.feature import MinMaxScaler
4 minMax = MinMaxScaler().setMin(5).setMax(10).setInputCol("features")
5 minMax.fit(scaledDF).transform(scaledDF).show(5, False)
```

id	features	MinMaxScaler_e07d77613bde__output
0	[1.0, 0.1, -1.0]	[5.0, 5.0, 5.0]
1	[2.0, 1.1, 1.0]	[7.5, 5.5, 7.5]
0	[1.0, 0.1, -1.0]	[5.0, 5.0, 5.0]
1	[2.0, 1.1, 1.0]	[7.5, 5.5, 7.5]
1	[3.0, 10.1, 3.0]	[10.0, 10.0, 10.0]

In [22]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import MaxAbsScaler
4 maScaler = MaxAbsScaler().setInputCol("features")
5 maScaler.fit(scaleDF).transform(scaleDF).show(5, False)
```

id	features	MaxAbsScaler_585c91198f51__output
0	[1.0, 0.1, -1.0]	[0.3333333333333333, 0.009900990099009903, -0.3333333333333333]
1	[2.0, 1.1, 1.0]	[0.6666666666666666, 0.10891089108910892, 0.3333333333333333]
0	[1.0, 0.1, -1.0]	[0.3333333333333333, 0.009900990099009903, -0.3333333333333333]
1	[2.0, 1.1, 1.0]	[0.6666666666666666, 0.10891089108910892, 0.3333333333333333]
1	[3.0, 10.1, 3.0]	[1.0, 1.0, 1.0]

In [23]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import ElementwiseProduct
4 from pyspark.ml.linalg import Vectors
5 scaleUpVec = Vectors.dense(10.0, 15.0, 20.0)
6 scalingUp = ElementwiseProduct()\
7     .setScalingVec(scaleUpVec)\
8     .setInputCol("features")
9 scalingUp.transform(scaleDF).show(5, False)
```

id	features	ElementwiseProduct_8d94b9364b23__output
0	[1.0, 0.1, -1.0]	[10.0, 1.5, -20.0]
1	[2.0, 1.1, 1.0]	[20.0, 16.5, 20.0]
0	[1.0, 0.1, -1.0]	[10.0, 1.5, -20.0]
1	[2.0, 1.1, 1.0]	[20.0, 16.5, 20.0]
1	[3.0, 10.1, 3.0]	[30.0, 151.5, 60.0]

In [24]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import Normalizer
4 manhattanDistance = Normalizer().setP(1).setInputCol("features")
5 manhattanDistance.transform(scaleDF).show(5, False)
```

```
+---+-----+-----+-----+
|id | features          | Normalizer_5f82750e3449__output |
+---+-----+-----+-----+
|0  | [1.0,0.1,-1.0] | [0.47619047619047616,0.047619047619047616,-0.47619047619047616] |
|1  | [2.0,1.1,1.0]  | [0.48780487804878053,0.26829268292682934,0.24390243902439027] |
|0  | [1.0,0.1,-1.0] | [0.47619047619047616,0.047619047619047616,-0.47619047619047616] |
|1  | [2.0,1.1,1.0]  | [0.48780487804878053,0.26829268292682934,0.24390243902439027] |
|1  | [3.0,10.1,3.0] | [0.18633540372670807,0.6273291925465838,0.18633540372670807] |
+---+-----+-----+-----+
```

In [26]:

```
1 # StringIndexer - transformer to convert categorical data into number
2
3 from pyspark.ml.feature import StringIndexer
4 lblIndxr = StringIndexer().setInputCol("lab").setOutputCol("labelInd")
5 idxRes = lblIndxr.fit(simpleDF).transform(simpleDF)
6 idxRes.show(5, False)
```

```
+-----+-----+-----+-----+-----+
|color|lab |value1|value2          |labelInd|
+-----+-----+-----+-----+-----+
|green|good|1      |14.386294994851129|1.0      |
|blue |bad |8      |14.386294994851129|0.0      |
|blue |bad |12     |14.386294994851129|0.0      |
|green|good|15     |38.97187133755819 |1.0      |
|green|good|12     |14.386294994851129|1.0      |
+-----+-----+-----+-----+-----+
```

only showing top 5 rows

```
In [27]: 1 # IndexToString - transformer to convert number back to category
2
3 from pyspark.ml.feature import IndexToString
4 labelReverse = IndexToString().setInputCol("labelInd")
5 labelReverse.transform(idxRes).show(5, False)
```

color	lab	value1	value2	labelInd	IndexToString_1e66c232baf4__output
green	good	1	14.386294994851129	1.0	good
blue	bad	8	14.386294994851129	0.0	bad
blue	bad	12	14.386294994851129	0.0	bad
green	good	15	38.97187133755819	1.0	good
green	good	12	14.386294994851129	1.0	good

only showing top 5 rows

```
In [28]: 1 # COMMAND -----
2
3 valIndexer = StringIndexer().setInputCol("value1").setOutputCol("value1_Ind")
4 valIndexer.fit(simpleDF).transform(simpleDF).show(5, False)
```

color	lab	value1	value2	value1_Ind
green	good	1	14.386294994851129	0.0
blue	bad	8	14.386294994851129	7.0
blue	bad	12	14.386294994851129	1.0
green	good	15	38.97187133755819	3.0
green	good	12	14.386294994851129	1.0

only showing top 5 rows

In [29]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import VectorIndexer
4 from pyspark.ml.linalg import Vectors
5 idxIn = spark.createDataFrame([
6     (Vectors.dense(1, 2, 3),1),
7     (Vectors.dense(2, 5, 6),2),
8     (Vectors.dense(1, 8, 9),3)
9 ]).toDF("features", "label")
10 indxr = VectorIndexer()\
11     .setInputCol("features")\
12     .setOutputCol("idxed")\
13     .setMaxCategories(2)
14 indxr.fit(idxIn).transform(idxIn).show(5, False)
```

```
+-----+-----+-----+
|features      |label|idxed      |
+-----+-----+-----+
|[1.0,2.0,3.0]|1    |[0.0,2.0,3.0]|
|[2.0,5.0,6.0]|2    |[1.0,5.0,6.0]|
|[1.0,8.0,9.0]|3    |[0.0,8.0,9.0]|
+-----+-----+-----+
```

In [30]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import OneHotEncoder, StringIndexer
4 lblIndxr = StringIndexer().setInputCol("color").setOutputCol("colorInd")
5 colorLab = lblIndxr.fit(simpleDF).transform(simpleDF.select("color"))
6 ohe = OneHotEncoder().setInputCol("colorInd")
7 ohe.transform(colorLab).show(5, False)
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-30-64f90509c5ec> in <module>
      5 colorLab = lblIndxr.fit(simpleDF).transform(simpleDF.select("color"))
      6 ohe = OneHotEncoder().setInputCol("colorInd")
----> 7 ohe.transform(colorLab).show(5, False)
```

AttributeError: 'OneHotEncoder' object has no attribute 'transform'

```
In [31]: 1 # Tokenizer - transformer to turn sentence into word array
2
3 from pyspark.ml.feature import Tokenizer
4 tkn = Tokenizer().setInputCol("Description").setOutputCol("DescOut")
5 tokenized = tkn.transform(sales.select("Description"))
6 tokenized.show(20, False)
```

Description	DescOut
RABBIT NIGHT LIGHT	[rabbit, night, light]
DOUGHNUT LIP GLOSS	[doughnut, lip, gloss]
12 MESSAGE CARDS WITH ENVELOPES	[12, message, cards, with, envelopes]
BLUE HARMONICA IN BOX	[blue, harmonica, in, box]
GUMBALL COAT RACK	[gumball, coat, rack]
SKULLS WATER TRANSFER TATTOOS	[skulls, , water, transfer, tattoos]
FELTCRAFT GIRL AMELIE KIT	[feltcraft, girl, amelie, kit]
CAMOUFLAGE LED TORCH	[camouflage, led, torch]
WHITE SKULL HOT WATER BOTTLE	[white, skull, hot, water, bottle]
ENGLISH ROSE HOT WATER BOTTLE	[english, rose, hot, water, bottle]
HOT WATER BOTTLE KEEP CALM	[hot, water, bottle, keep, calm]
SCOTTIE DOG HOT WATER BOTTLE	[scottie, dog, hot, water, bottle]
ROSE CARAVAN DOORSTOP	[rose, caravan, doorstop]
GINGHAM HEART DOORSTOP RED	[gingham, heart, , doorstop, red]
STORAGE TIN VINTAGE LEAF	[storage, tin, vintage, leaf]
SET OF 4 KNICK KNACK TINS POPPIES	[set, of, 4, knick, knack, tins, poppies]
POPCORN HOLDER	[popcorn, holder]
GROW A FLYTRAP OR SUNFLOWER IN TIN	[grow, a, flytrap, or, sunflower, in, tin]
AIRLINE BAG VINTAGE WORLD CHAMPION	[airline, bag, vintage, world, champion]
AIRLINE BAG VINTAGE JET SET BROWN	[airline, bag, vintage, jet, set, brown]

only showing top 20 rows

In [33]:

```
1 # RegexTokenizer - Tokenizer with RegEx
2
3 from pyspark.ml.feature import RegexTokenizer
4 rt = RegexTokenizer()\
5     .setInputCol("Description")\
6     .setOutputCol("DescOut")\
7     .setPattern(" ") \
8     .setToLowercase(True)
9 rt.transform(sales.select("Description")).show(10, False)
```

Description	DescOut
RABBIT NIGHT LIGHT	[rabbit, night, light]
DOUGHNUT LIP GLOSS	[doughnut, lip, gloss]
12 MESSAGE CARDS WITH ENVELOPES	[12, message, cards, with, envelopes]
BLUE HARMONICA IN BOX	[blue, harmonica, in, box]
GUMBALL COAT RACK	[gumball, coat, rack]
SKULLS WATER TRANSFER TATTOOS	[skulls, water, transfer, tattoos]
FELTCRAFT GIRL AMELIE KIT	[feltcraft, girl, amelie, kit]
CAMOUFLAGE LED TORCH	[camouflage, led, torch]
WHITE SKULL HOT WATER BOTTLE	[white, skull, hot, water, bottle]
ENGLISH ROSE HOT WATER BOTTLE	[english, rose, hot, water, bottle]

only showing top 10 rows

In [34]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import RegexTokenizer
4 rt = RegexTokenizer()\
5     .setInputCol("Description")\
6     .setOutputCol("DescOut")\
7     .setPattern(" ")\
8     .setGaps(False)\
9     .setToLowercase(True)
10 rt.transform(sales.select("Description")).show(10, False)
```

```
+-----+-----+
|Description                |DescOut          |
+-----+-----+
|RABBIT NIGHT LIGHT         |[ , ]            |
|DOUGHNUT LIP GLOSS         |[ , , ]          |
|12 MESSAGE CARDS WITH ENVELOPES|[ , , , ]        |
|BLUE HARMONICA IN BOX      |[ , , , ]        |
|GUMBALL COAT RACK          |[ , ]            |
|SKULLS WATER TRANSFER TATTOOS|[ , , , , ]      |
|FELTCRAFT GIRL AMELIE KIT  |[ , , ]          |
|CAMOUFLAGE LED TORCH       |[ , ]            |
|WHITE SKULL HOT WATER BOTTLE|[ , , , , ]      |
|ENGLISH ROSE HOT WATER BOTTLE|[ , , , ]        |
+-----+-----+
```

only showing top 10 rows

In [35]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import StopWordsRemover
4 englishStopWords = StopWordsRemover.loadDefaultStopWords("english")
5 stops = StopWordsRemover()\
6     .setStopWords(englishStopWords)\
7     .setInputCol("DescOut")
8 stops.transform(tokenized).show(5, False)
```

```
+-----+-----+-----+
+-----+
|Description                |DescOut                |StopWordsRemover_352087335d7
d__output|
+-----+-----+-----+
+-----+
|RABBIT NIGHT LIGHT        |[rabbit, night, light] |[rabbit, night, light]
|
|DOUGHNUT LIP GLOSS        |[doughnut, lip, gloss] |[doughnut, lip, gloss]
|
|12 MESSAGE CARDS WITH ENVELOPES|[12, message, cards, with, envelopes]|[12, message, cards, envelop
es]
|
|BLUE HARMONICA IN BOX      |[blue, harmonica, in, box] |[blue, harmonica, box]
|
|GUMBALL COAT RACK          |[gumball, coat, rack] |[gumball, coat, rack]
|
+-----+-----+-----+
+-----+
only showing top 5 rows
```

In [36]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import NGram
4 unigram = NGram().setInputCol("DescOut").setN(1)
5 bigram = NGram().setInputCol("DescOut").setN(2)
6 unigram.transform(tokenized.select("DescOut")).show(5, False)
7 bigram.transform(tokenized.select("DescOut")).show(5, False)
```

```
+-----+-----+
|DescOut|NGram_89b73e4d33d9__output|
+-----+-----+
|[rabbit, night, light]|
|[doughnut, lip, gloss]|
|[12, message, cards, with, envelopes]|
|[blue, harmonica, in, box]|
|[gumball, coat, rack]|
+-----+-----+
only showing top 5 rows
```

```
+-----+-----+
|DescOut|NGram_09d12ca8503c__output|
+-----+-----+
|[rabbit, night, light]|
|[doughnut, lip, gloss]|
|[12, message, cards, with, envelopes]|
|[blue, harmonica, in, box]|
|[gumball, coat, rack]|
+-----+-----+
only showing top 5 rows
```

In [37]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import CountVectorizer
4 cv = CountVectorizer()\
5     .setInputCol("DescOut")\
6     .setOutputCol("countVec")\
7     .setVocabSize(500)\
8     .setMinTF(1)\
9     .setMinDF(2)
10 fittedCV = cv.fit(tokenized)
11 fittedCV.transform(tokenized).show(5, False)
```

```
+-----+-----+-----+
+-----+
|Description                |DescOut                |countVec
|
+-----+-----+-----+
+-----+
|RABBIT NIGHT LIGHT        |[rabbit, night, light] |(500,[150,185,212],[1.0,1.0,
1.0]) |
|DOUGHNUT LIP GLOSS        |[doughnut, lip, gloss] |(500,[462,463,491],[1.0,1.0,
1.0]) |
|12 MESSAGE CARDS WITH ENVELOPES|[12, message, cards, with, envelopes]|(500,[35,41,166],[1.0,1.0,1.
0]) |
|BLUE HARMONICA IN BOX      |[blue, harmonica, in, box] |(500,[10,16,36,352],[1.0,1.
0,1.0,1.0])|
|GUMBALL COAT RACK          |[gumball, coat, rack] |(500,[228,281,407],[1.0,1.0,
1.0]) |
+-----+-----+-----+
```

only showing top 5 rows

```
In [38]: 1 # COMMAND -----
2
3 tfIdfIn = tokenized\
4   .where("array_contains(DescOut, 'red')")\
5   .select("DescOut")\
6   .limit(10)
7 tfIdfIn.show(10, False)
```

```
+-----+
|DescOut|
+-----+
|[gingham, heart, , doorstop, red]|
|[red, floral, feltcraft, shoulder, bag]|
|[alarm, clock, bakelike, red]|
|[pin, cushion, babushka, red]|
|[red, retrospot, mini, cases]|
|[red, kitchen, scales]|
|[gingham, heart, , doorstop, red]|
|[large, red, babushka, notebook]|
|[red, retrospot, oven, glove]|
|[red, retrospot, plate]|
+-----+
```

```
In [39]: 1 # COMMAND -----
2
3 from pyspark.ml.feature import HashingTF, IDF
4 tf = HashingTF()\
5   .setInputCol("DescOut")\
6   .setOutputCol("TFOut")\
7   .setNumFeatures(10000)
8 idf = IDF()\
9   .setInputCol("TFOut")\
10  .setOutputCol("IDFOut")\
11  .setMinDocFreq(2)
```

In [40]:

```
1 # COMMAND -----
2
3 idf.fit(tf.transform(tfIdfIn)).transform(tf.transform(tfIdfIn)).show(10, False)
```

```
+-----+-----+-----+
-----+
|DescOut          |TFOut          |ID
FOut
|
+-----+-----+-----+
-----+
|[gingham, heart, , doorstop, red]      |(10000,[52,804,3372,6594,9808],[1.0,1.0,1.0,1.0,1.0])|(1
0000,[52,804,3372,6594,9808],[0.0,1.2992829841302609,1.2992829841302609,1.2992829841302609,1.2992
829841302609])|
|[red, floral, feltcraft, shoulder, bag]|(10000,[50,52,415,6756,8005],[1.0,1.0,1.0,1.0,1.0]) |(1
0000,[50,52,415,6756,8005],[0.0,0.0,0.0,0.0,0.0])
|
|[alarm, clock, bakelike, red]           |(10000,[52,4995,8737,9001],[1.0,1.0,1.0,1.0]) |(1
0000,[52,4995,8737,9001],[0.0,0.0,0.0,0.0])
|
|[pin, cushion, babushka, red]           |(10000,[52,610,2490,7153],[1.0,1.0,1.0,1.0]) |(1
0000,[52,610,2490,7153],[0.0,0.0,0.0,1.2992829841302609])
|
|[red, retrospot, mini, cases]           |(10000,[52,547,6703,8448],[1.0,1.0,1.0,1.0]) |(1
0000,[52,547,6703,8448],[0.0,0.0,0.0,1.0116009116784799])
|
|[red, kitchen, scales]                 |(10000,[52,756,6452],[1.0,1.0,1.0]) |(1
0000,[52,756,6452],[0.0,0.0,0.0])
|
|[gingham, heart, , doorstop, red]      |(10000,[52,804,3372,6594,9808],[1.0,1.0,1.0,1.0,1.0])|(1
0000,[52,804,3372,6594,9808],[0.0,1.2992829841302609,1.2992829841302609,1.2992829841302609,1.2992
829841302609])|
|[large, red, babushka, notebook]        |(10000,[52,2787,7022,7153],[1.0,1.0,1.0,1.0]) |(1
0000,[52,2787,7022,7153],[0.0,0.0,0.0,1.2992829841302609])
|
|[red, retrospot, oven, glove]           |(10000,[52,8242,8448,8667],[1.0,1.0,1.0,1.0]) |(1
0000,[52,8242,8448,8667],[0.0,0.0,1.0116009116784799,0.0])
|
|[red, retrospot, plate]                 |(10000,[52,4925,8448],[1.0,1.0,1.0]) |(1
0000,[52,4925,8448],[0.0,0.0,1.0116009116784799])
```

```
|
+-----+-----+
-----
-----+
-----+
```

In [42]: 1 display(idf.fit(tf.transform(tfIdfIn)).transform(tf.transform(tfIdfIn)).limit(10).toPandas())

	DescOut	TFOut	IDFOut
0	[gingham, heart, , doorstep, red]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
1	[red, floral, feltcraft, shoulder, bag]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
2	[alarm, clock, bakelike, red]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
3	[pin, cushion, babushka, red]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
4	[red, retrospot, mini, cases]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
5	[red, kitchen, scales]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
6	[gingham, heart, , doorstep, red]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
7	[large, red, babushka, notebook]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
8	[red, retrospot, oven, glove]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
9	[red, retrospot, plate]	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...	(0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...

In [43]:

```
1  # COMMAND -----
2
3  from pyspark.ml.feature import Word2Vec
4  # Input data: Each row is a bag of words from a sentence or document.
5  documentDF = spark.createDataFrame([
6      ("Hi I heard about Spark".split(" "), ),
7      ("I wish Java could use case classes".split(" "), ),
8      ("Logistic regression models are neat".split(" "), )
9  ], ["text"])
10 # Learn a mapping from words to Vectors.
11 word2Vec = Word2Vec(vectorSize=3, minCount=0, inputCol="text",
12     outputCol="result")
13 model = word2Vec.fit(documentDF)
14 result = model.transform(documentDF)
15 for row in result.collect():
16     text, vector = row
17     print("Text: [%s] => \nVector: %s\n" % (" ".join(text), str(vector)))
```

Text: [Hi, I, heard, about, Spark] =>

Vector: [0.10983876287937165, -0.03447718722745776, 0.000940057821571827]

Text: [I, wish, Java, could, use, case, classes] =>

Vector: [0.0072656965681484765, -0.018058971102748598, 0.003386378288269043]

Text: [Logistic, regression, models, are, neat] =>

Vector: [-0.047607143968343736, -0.0320490337908268, 0.07224417026154697]

In [44]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import PCA
4 pca = PCA().setInputCol("features").setK(2)
5 pca.fit(scaleDF).transform(scaleDF).show(20, False)
```

id	features	PCA_6f4ed9b6b39d__output
0	[1.0, 0.1, -1.0]	[0.0713719499248417, -0.4526654888147805]
1	[2.0, 1.1, 1.0]	[-1.6804946984073723, 1.2593401322219198]
0	[1.0, 0.1, -1.0]	[0.0713719499248417, -0.4526654888147805]
1	[2.0, 1.1, 1.0]	[-1.6804946984073723, 1.2593401322219198]
1	[3.0, 10.1, 3.0]	[-10.872398139848944, 0.030962697060155975]

In [45]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import PolynomialExpansion
4 pe = PolynomialExpansion().setInputCol("features").setDegree(2)
5 pe.transform(scaleDF).show(5, False)
```

```
+---+-----+-----+
----+
|id|features      |PolynomialExpansion_70c2c7232672__output|
|
+---+-----+-----+
----+
|0| [[1.0,0.1,-1.0]]|[1.0,1.0,0.1,0.1,0.010000000000000002,-1.0,-1.0,-0.1,1.0]|
|
|1| [[2.0,1.1,1.0]]|[2.0,4.0,1.1,2.2,1.2100000000000002,1.0,2.0,1.1,1.0]|
|
|0| [[1.0,0.1,-1.0]]|[1.0,1.0,0.1,0.1,0.010000000000000002,-1.0,-1.0,-0.1,1.0]|
|
|1| [[2.0,1.1,1.0]]|[2.0,4.0,1.1,2.2,1.2100000000000002,1.0,2.0,1.1,1.0]|
|
|1| [[3.0,10.1,3.0]]|[3.0,9.0,10.1,30.299999999999997,102.00999999999999,3.0,9.0,30.299999999999997,9.0]|
+---+-----+-----+
----+
```

```
In [46]: 1 # COMMAND -----
2
3 from pyspark.ml.feature import ChiSqSelector, Tokenizer
4 tkn = Tokenizer().setInputCol("Description").setOutputCol("DescOut")
5 tokenized = tkn\
6     .transform(sales.select("Description", "CustomerId"))\
7     .where("CustomerId IS NOT NULL")
8 prechi = fittedCV.transform(tokenized)\
9     .where("CustomerId IS NOT NULL")
10 chisq = ChiSqSelector()\
11     .setFeaturesCol("countVec")\
12     .setLabelCol("CustomerId")\
13     .setNumTopFeatures(2)
14 chisq.fit(prechi).transform(prechi)\
15     .drop("customerId", "Description", "DescOut").show(5, False)
```

countVec	ChiSqSelector_9631cfd8c94f__output
(500, [150, 185, 212], [1.0, 1.0, 1.0])	(2, [], [])
(500, [462, 463, 491], [1.0, 1.0, 1.0])	(2, [], [])
(500, [35, 41, 166], [1.0, 1.0, 1.0])	(2, [], [])
(500, [10, 16, 36, 352], [1.0, 1.0, 1.0, 1.0])	(2, [], [])
(500, [228, 281, 407], [1.0, 1.0, 1.0])	(2, [], [])

only showing top 5 rows

```
In [47]: 1 # COMMAND -----
2
3 fittedPCA = pca.fit(scaleDF)
4 fittedPCA.write().overwrite().save("/tmp/fittedPCA")
```

In [48]:

```
1 # COMMAND -----
2
3 from pyspark.ml.feature import PCAModel
4 loadedPCA = PCAModel.load("/tmp/fittedPCA")
5 loadedPCA.transform(scaledDF).show(5, False)
6
7
8 # COMMAND -----
```

```
+---+-----+-----+
|id|features          |PCAModel_5f0cec389764__output|
+---+-----+-----+
|0| |[1.0,0.1,-1.0]| |[0.0713719499248417,-0.4526654888147805]|
|1| |[2.0,1.1,1.0]| |[-1.6804946984073723,1.2593401322219198]|
|0| |[1.0,0.1,-1.0]| |[0.0713719499248417,-0.4526654888147805]|
|1| |[2.0,1.1,1.0]| |[-1.6804946984073723,1.2593401322219198]|
|1| |[3.0,10.1,3.0]| |[-10.872398139848944,0.030962697060155975]|
+---+-----+-----+
```

In [49]:

```
1 spark.stop()
```

In []:

```
1
```